

# Memory Isolation as a Kernel Primitive in Dune

Adrien Ghosn  
*EPFL*

Second Name  
*Second Institution*

## Abstract

We believe programmers would greatly benefit from an OS abstraction for memory isolation within a process. We further argue that relying on Dune and hardware support for virtualization would, unlike previous attempts, yield a maintainable solution. Dune provides flexibility, compatibility with the Linux kernel, and access to hardware features required to implement such an abstraction. We improve upon previous solutions by both relying on hardware mechanisms to enforce memory isolation and being completely decoupled from an existing kernel.

## 1 Introduction

General purpose kernels rely on the process abstraction to represent a unit of isolation in terms of memory, privilege, and execution. Processes isolate applications from each other and, more specifically, prevent one process from accessing another's memory. Although vital for operating systems, this abstraction requires process management, context switching, and resource accountability which, in turn, introduce non-negligible overheads at the OS-level, especially compared to the associated hardware costs.

We identify a discrepancy between the functionalities provided by existing general purpose kernels and user needs. While the thread abstraction provides independent, isolated, execution units within a process, memory isolation is attainable only by going through the costly creation and management of a new process. In fact, the lack of a kernel primitive to isolate memory units within the same thread of execution, pushed application developers to implement their own solutions within user space.

Memory isolation within a single thread of execution is a generalization of the ring protection mechanism,

a common technique used in general purpose kernel to isolate kernel and user spaces. Mode switching is a technique that allows to perform privileged operations without having to incur the cost of a context switch. Kernel and user space isolation is enforced by efficient hardware mechanisms. For example, page table entries that belong to the kernel are marked with the supervisor bit, hence putting them out of reach of user applications.

Programming languages, and more generally user applications, have attempted to provide a memory isolation abstraction within a single thread of execution. For example, some object oriented languages like Java, allow to mark class attributes as private, i.e., the object cannot be accessed from some parts of the program. This isolation is enforced using compile and runtime checks. The common technique is therefore to rely on a user application to enforce user to user isolation, i.e., to limit access to certain memory regions. This is both inefficient, since it requires software interposition on memory accesses, and hardly implemented. User applications cannot rely on hardware to enforce isolation and, more often than not, present vulnerabilities and ways to circumvent isolation mechanisms. For example, in Java, meta-programming allows to access even the private fields.

We believe that memory isolation within a single thread of execution must be provided as a kernel primitive. Furthermore, we argue that, unlike previous implementations, it is possible to provide a maintainable and flexible solution, easy to tune to special use cases, by relying on efficient virtualization and direct access to hardware features. We implement the existing lightweight contexts abstraction[REF] (lwC) as a Dune library, and argue about the advantages of having such a solution within Dune[REF]. This paper is organized as follows: Section[REF] presents previous implementations that strove to provide a memory isolation abstrac-

tion, and their limitations. Section [REF] presents the *lwC* API, as exposed in our Dune implementation. Section [REF] presents how Dune was modified to allow the implementation of the *lwC* abstraction. Finally, Section [REF] lists future work before the conclusion in Section [REF].

## 2 Shortcomings of Existing Solutions

Computer scientists have made several attempts at providing a memory isolation abstraction within a single thread of execution. We observed two distinct techniques to provide this abstraction: sandboxes that isolate different user memory regions, and kernel level implementations. In this section, we briefly describe one sandbox solution, and explain why we believe that this abstraction has to be provided as a kernel primitive. We then proceed with existing kernel level implementations, and expose what we consider to be important limitations that prevent them from being a stable long-term solution.

A sandbox is a security mechanism that allows to isolate untrusted code execution and prevent it from harming both the underlying system and other sandboxed applications. A common requirement for sandboxes is to enforce memory isolation and reduce the capabilities of running applications. As a result, sandboxes often have to intervene in the management of resources. They hence duplicate the kernel's functionalities in order to obtain fine-grained control over memory allocation and accesses. Google's NaCl [REF] provides a sandboxed environment to execute untrusted native plugins in a Web browser. In NaCl, memory isolation is ensured using code instrumentation. The untrusted application's code is recompiled and instrumented such that instructions that access and modify memory related registers are replaced by guarded pseudo instructions, hence insuring that the registers content are consistent throughout the execution. All other memory accesses are re-written to be expressed using these guarded registers. NaCl can therefore ensure that every memory operation falls into the proper address range and no untrusted application can access another's resources. NaCl authors claim that control-flow and memory integrity are insured with an average overhead of less than 7% on x86-64. While these performances are remarkable for a solution based on code instrumentation, we believe that the recompilation step required to perform the binary translation, as well as the introduction of pseudo-instructions are a burden that could be avoided if hardware support for memory isolation was used. As it is the kernel's role to provide resource management and isolation, we argue that this abstraction should be provided as a collection of system calls.

Kernels are at the core of operating systems and are responsible for managing resources, among which the memory. As such, they are a perfect environment to implement and expose new abstractions for memory isolation within a process. Not only does it fit in the kernel's role, but also allows to rely on hardware features to enforce memory isolation. Others [REF] before us identified the need for a new memory isolation abstraction, and the advantages of implementing it as a kernel primitive. The following two examples both rely on page tables and hardware features to implement their abstractions.

Wedge [REF] is a kernel library that introduces the *sthread* abstraction. *Sthreads* are described as threads with fork-like semantics. More specifically, upon creation, modifiers allow a fine-grained control over resource visibility within the newly created *sthread*. Entire memory regions can be unmapped and made inaccessible from within the *sthread*. This solution, however, does not completely decouple the memory and execution units. *Sthreads* are schedulable entities and therefore provide memory isolation at a cost that comprises *sthreads* scheduling and management. They therefore can be seen as a slightly cheaper version process abstraction.

Light-weight contexts [REF], *lwCs* for short, are implemented within the FreeBSD kernel and expose the *context*, or *lwC*, abstraction. Much like *sthreads*, users have a fine-grained control over memory mappings when creating a *context*. Entire address ranges can be shared, copied, unmapped, are made read-only. Unlike *sthreads*, however, *contexts* are orthogonal to threads, i.e., they are non-schedulable entities. A single thread of execution can switch between different *contexts*. Upon a switch, the execution state is saved within the current *lwC*. This solution completely decouples the memory and execution isolation mechanisms. However, the authors had to modify the FreeBSD kernel in order to implement *lwCs*. This new feature impacts memory management and might therefore introduce new vulnerabilities in the kernel due to software bugs. As a result, one should not expect *lwCs* to be merged back inside the FreeBSD's master branch. We can therefore wonder how maintainable this solution really is. In fact, extra work will be required to merge back bug fixes and updates made to the original kernel into the *lwC* branch, and thus raising the question of long-time support for the *lwC* library. Without long-time support, we doubt that *lwCs* will be adopted and used as a building block for new projects. More importantly, implementing a solution within an existing kernel requires to interact with rigid APIs and re-

spect the kernel's implementation. From an implementation point of view, this reduces the solution's flexibility and hence might rule-out some particular implementations that would prove to be more efficient in terms of performances. For example, the *lwCs* authors had to work with FreeBSD's datastructures and abstraction layers for memory management. We argue that, in a virtualized environment with hardware support, we can implement a similar solution, with equivalent performances, and more flexibility and freedom in terms of the memory hierarchy abstractions and the actual implementation of *lwC* functionalities. Moreover, our solution does not require to modify an existing kernel.

### 3 Leveraging hardware support for virtualization

Intel and AMD identified and answered the need for virtualization hardware support. The Intel solution, called VT-x, is a virtualization extension to the x86 ISA. AMD provides a similar extension with SVM[REF]. Instead of introducing traps for each instruction accessing privileged state and relying on a VMM for emulation, these extensions strive to execute as many instructions, privileged or not, directly in hardware. To avoid involving the VMM and hence introduce delays, the hardware maintains a shadow copy of privileged state. In Intel VT-x, the CPU provides two modes of operation: VMX root and VMX non-root. Guest OSes, i.e., virtualized environments, run in VMX non-root mode where CPU behavior is restricted. Transitions between the two modes, e.g., when the VMM needs to be involved, is managed by hardware. Transitions from non-root to root and from root to non-root are called *VM exits* and *VM entries* respectively.

Hardware support for virtualization provides direct access to hardware features, such as ring protection mechanism, page tables, TLBs, from within a virtualized environment, i.e., the guest. Such features are essential to implement and enforce efficient memory isolation within the same thread of execution, as discussed previously. Using hardware support for virtualization, we therefore have an environment that is, on one hand, independent from the host, i.e., the original kernel, and, at the same time, undistinguishable from it in terms of hardware tools made available.

While a certain form of hypervisor is still needed to interface the guest and host environments, we argue that hardware support for virtualization, as much as our own limited requirements, enable to reduce the hypervisor functionalities and complexity. By design, a hypervisor

with a narrow set of features is less amenable to exploits than a standard one. Any hypervisor can be reduced to its core: the exposed functionalities, and its management of resources. Where standard hypervisors such as[CITE] accumulate features and functionalities, emulating a complete operating system, we believe that our needs, in order to provide memory isolation within a thread of execution can be strip down to a bare minimum consisting of an execution unit and some memory. Our hypervisor will therefore be solely in charge of exposing a certain amount of memory resources, i.e., virtualizing memory resources. All other functionalities are implemented within the guest, by relying on direct access to hardware features and privileged state.

### 4 Adapting Dune

TODO:Memory management must be modified. Multiple address spaces must be allowed + less rigid mapping with host. Interposition on syscalls might be consider to limit access? Have to do all of that without introducing significant delays Would like to preserve backward compatibility with existing Dune applications. Once all of that is proven, we can merge back into Dune main branch. Easier to commit to Dune than to Linux ;)

### 5 Future Work

TODO: performance evaluation. Re-implement existing applications using *lwc*.

### 6 Conclusion

### 7 Acknowledgments

### References

### Notes

<sup>1</sup>Remember to use endnotes, not footnotes!